WEST UNIVERSITY OF TIMIŞOARA
FACULTY OF MATHEMATICS AND COMPUTER
SCIENCE
MASTER STUDY PROGRAM: SOFTWARE
ENGINEERING

Verification of ACAS-Xu benchmark using Alpha Beta Crown and Nnenum

Supervisor:
Mădălina Erașcu

Students:
Amalia Duma Diana
Răzvan Maciovan Alexandru
Luca Mitroi
Vlad Petcu
Vlad Țîrcomnicu

TIMIŞOARA
2024

Abstract

This paper evaluates the effectiveness of the Alpha Beta Crown and Nnenum verification tools with the ACAS-Xu benchmark, which is critical for ensuring the safety of unmanned aircraft systems by issuing turn rate advisories. The project aimed to replicate findings from the VNN-COMP 2023 competition, focusing on the verification of neural networks for the ACAS-Xu system. The dataset uses parameters from sensor measurements across 45 deep neural networks, with the verification targeting 10 specific properties. The results indicated that out of 186 instances, 139 were safe, and 47 unsafe, achieving an accuracy of approximately 74.73%. Besides the results, the study showcases the challenges of the tools setup.

1   Introduction

The assignment for this project, required the usage of one benchmark and two tools for which to analyze upon. For the purpose of our project, we have decided to utilize the ACAS-Xu benchmark[1], along with the tools Alpha Beta Crown[2] and Nnenum [3]. In the following sections we will see what the selected benchmark represents (section 2), the challenges the tools presented (section 3), as well as how ACAS-Xu scored with the chosen tools based on the VNN COMP 2023 competition (section 4).

2   ACAS-Xu Dataset

ACAS stands for Airborne Collision Avoidance System. There are multiple types of ACAS benchmarks, but the one this paper is based on is Xu, which is optimized for unmanned aircraft systems (UAS), issuing turn rate advisories to remote pilots[3]. The installation process for the benchmark was easy, only needing to download it from GitHub[1].

The input data that can be found across the VNN-LIB files, revolves around different aircraft parameters determined from sensor measurements[5]. These parameters can be seen in figure 1:

- ρ - Distance from ownership to intruder;

- θ - Angle to intruder relative to aircraft heading direction length;

- ψ - Heading angle of intruder relative to aircraft heading direction;

- v_{own} - Speed of aircraft;

- v_{int} - Speed of intruder;

- τ - Time until loss of vertical separation;

- a_{prev} - Previous advisory;

These values were introduced into the neural network (the ONNX files), representing a combination of 45 deep neural networks. The choice of having 45 DNNs comes from the need to reduce the lookup time, as previously there was only one DNN with 2GB in size [3]. They were created by mixing the time until loss of vertical separation (τ) and the previous advisory (a_{prev}). Therefor every network has five inputs and five outputs. The inputs are the aforementioned ones (excluding τ and a_{prev}), while the outputs represent the different horizontal advisories provided to the aircraft: Clear-of-Conflict (COC), weak right, strong right, weak left, or strong left. The DNNs are fully connected, use ReLU activation functions, and have six hidden layers with a total of 300 ReLU nodes each. [3].

In order to confirm the correctness of the results, figure 2 illustrates a top-down of a head-on encounter scenario, where each pixel shows a different course of action based on the presence of an intruder.

---

[1]`https://github.com/ChristopherBrix/vnncomp2023_benchmarks/tree/main/benchmarks/acasxu`

[2]`https://github.com/Verified-Intelligence/alpha-beta-CROWN`

[3]`https://github.com/stanleybak/nnenum`

Figure 1: Geometry of ACAS Xu horizontal logic table [3]

Figure 2: Advisories for a head-on encounter with $a_{\text{prev}} = $ COC and $\tau = $ os [3]

3 Tools

For this verification assignment, two tools were used to ensure the reliability and safety of ACAS XU: alpha-beta-CROWN and nnenum. Alpha Beta Crown serves as a neural network verifier based on an efficient linear bound propagation framework, built on bound-propagation-based neural network verifiers: CROWN [8, 11], alpha-CROWN [9], beta-CROWN [7], GCP-CROWN [10], and nonlinear branch-and-bound [6]. Nnenum is a verifier that relies on multiple levels of abstraction to achieve high performance verification of ReLU networks without sacrificing completeness [1]. In the following sections, the setup of these tools is presented.

3.1 Alpha-Beta-CROWN

Setting it up involved a two-step process: first, installing Miniconda and then the tool itself. There were no complications in the first part, as there is an executable file available on their official site that takes care of the entire setup once run. However, the latter part, involving the tool's installation, presented numerous complications.

The setup of Alpha Beta Crown followed their "Installation and Setup" section outlined on GitHub. The first command was supposed to clone their GitHub repository including "auto_LiRPA"[8], an essential component for the tool's functionality. Despite running this command, only the tool's files were retrieved due to permission denials that restricted access to the second GitHub repository. Resolving this required an additional step: manually obtaining the missing necessary files and placing them in their indicated location. The next phase involved creating the conda environment using the specified command, which led to an error that can be observed in figure 3.

This error consumed a considerable amount of time, as the error message initially suggested an issue with the script. After extensive investigation, the root cause was traced to the configuration file. Instead of containing the necessary configuration for the conda environment, it referenced another configuration file by name. After resolving the configuration file issue, the installation of the tool seemed successful. However, when attempting to run it, another problem emerged indicating a missing library, "auto_LiRPA". Although the necessary files existed, the error message

4

```
# >>>>>>>>>>>>>>>>>>>>> ERROR REPORT <<<<<<<<<<<<<<<<<<<<<<
Traceback (most recent call last):
  File "/home/amalia/miniconda3/lib/python3.11/site-packages/conda_env/env.py", line 162, in from_yaml
    data = validate_keys(data, kwargs)
           ^^^^^^^^^^^^^^^^^^^^^^^^^^^^
  File "/home/amalia/miniconda3/lib/python3.11/site-packages/conda_env/env.py", line 30, in validate_keys
    new_data = data.copy() if data else {}
               ^^^^^^^^^^^
AttributeError: 'str' object has no attribute 'copy'
```

Figure 3: Creating conda environment error

revealed a discrepancy in the file path. It seemed that the correct location for this library was wrongly specified. With this adjustment, the setup concluded, and the tool run successfully.

3.2 Nnenum

As a prerequisite, the tool required Docker to be installed, which simplified the process by eliminating all dependencies challenges thanks to the docker container properties. The setup followed their "Getting started" section outlined in Github. For the actual installation of the tool the GitHub repository needs to be cloned. There are two commands that needs to be run in order for the tool to work. The first one is for building the image and the second one is for getting the shell after building the image. Both commands are in the docker file present in the tool repository.

If both commands run successfully the tool is installed and running, ready to be used. To double check if the tool is running the container should be visible in the docker desktop application, which can be observed in figure 4. The installation of the tool was simple and did not produce any error or unexpected behavior.

| | | sharp_merkle 6893e0883622 | nnenum_image | Running | 0% | 32 seconds ago | ■ ⋮ 🗑 |

Figure 4: Nnenum container running

4 Experimental Results

The verification of ACAS Xu benchmark involved evaluating the 10 properties against the 45 DNNs in the following manner: the first four were checked across all 45 networks, whereas the rest were checked on a single network. This approach followed the work done by the original authors [4]. For each property, the timeout was set to 116 seconds. This verification process was performed using alpha-beta-CROWN and nnenum. Initially, the testing focused solely on property 1 against the first network to gain insights into the output.

The output from alpha-beta-CROWN contains details about configuration, attack parameters, model output, properties, verification iterations, results, and a summary. Every iteration provides data about the batch size, the worst lower bound encountered, the total time spent in that iteration, the number of domains or branches explored, the cumulative time and the current lower bound - right-hand

side value. This particular value is used to asses whether the property is satisfied or not, since a negative value indicates a potential violation of the property. One example of such iteration can be observed in figure 5.

Figure 5: Alpha-Beta-CROWN first iteration in the verification process

In contrast to alpha-beta-CROWN, nnenum differs in that it produces more concise output, focusing primarily on the verification result and essential statistics such as runtime, work fraction completed, and the count of symbolic abstract states ('stars') evaluated. In the case of an unsafe result, a counterexample that demonstrates the violation is also shown. One such output is presented in figure 6.

Figure 6: nnenum unsafe output

The next phase involved testing all properties, aiming to evaluate the neural network's behavior across various configurations and constraints. The final verified accuracy stands at approximately 74.73%. Upon comparing the results from alpha-beta-CROWN and nnenum for the 186 instances tested, both tools identified 139 instances as safe, meaning that the property held true, and 47 instances as unsafe. This information is summarized in table 1.

Tool	Verified	Falsified	Penalty	Score
alpha-beta-CROWN	139	47	0	1860
nnenum	139	47	0	1860

Table 1: Benchmark 2023-ACAS-Xu

Regarding the verification times for alpha-beta-CROWN, the mean time taken for verifying all instances was approximately 2.58 seconds, with the maximum time being around 69.08 seconds. In terms of the employed methods, specifically the

Branch and Bound method (BAB), 4 instances were marked as unsafe. In contrast, utilizing the Projected Gradient Descent (PGD) method, 47 instances were identified as unsafe.

As for nnenum, the performance achieved was quite similar, where all the instances have been verified, on average, in approximately 2.3 seconds with a maximum verification time of 61.58.

A more detailed overview is available on GitHub[4], featuring two CSV files that present results for both alpha-beta-CROWN and nnenum. These files detail the satisfiability and verification times for every tested combination. In addition, counter examples were generated and uploaded for the instances marked as unsafe.

5   Conclusions

This study presents the results of the evaluation of state of the art verification methods on the ACAS-Xu benchmark. The aim was to replicate the results from the VNN-COMP 2023[2] competition with the usage of alpha-beta-CROWN and nnenum. The findings revealed that out of 186 instances examined, 47 were identified as unsafe. It's worth noting a significant difference in the results when running alpha-beta-CROWN compared to the competition, as there were no penalties. The instance that did result in a penalty was classified as unsafe. The verification process used the same parameters as those used in the competition, with the source filed obtained directly from the official GitHub repository [5].

---

[4]`https://github.com/razvanmaciovan/Proiect-VF`
[5]`https://github.com/ChristopherBrix/vnncomp2023_benchmarks/tree/main/`
`benchmarks/acasxu`

References

[1] Stanley Bak. nnenum: Verification of relu neural networks with optimized abstraction refinement. In *NASA Formal Methods: 13th International Symposium, NFM 2021, Virtual Event, May 24–28, 2021, Proceedings*, page 19–36, Berlin, Heidelberg, 2021. Springer-Verlag.

[2] Christopher Brix, Stanley Bak, Changliu Liu, and Taylor T. Johnson. The fourth international verification of neural networks competition (vnn-comp 2023): Summary and results, 2023.

[3] Guy Katz, Clark Barrett, David L Dill, Kyle Julian, and Mykel J Kochenderfer. Reluplex: An efficient SMT solver for verifying deep neural networks. In *Computer Aided Verification: 29th International Conference, CAV 2017, Heidelberg, Germany, July 24-28, 2017, Proceedings, Part I 30*, pages 97–117. Springer, 2017.

[4] Guy Katz, Derek A. Huang, Duligur Ibeling, Kyle Julian, Christopher Lazarus, Rachel Lim, Parth Shah, Shantanu Thakoor, Haoze Wu, Aleksandar Zeljić, David L. Dill, Mykel J. Kochenderfer, and Clark Barrett. The marabou framework for verification and analysis of deep neural networks. In Isil Dillig and Serdar Tasiran, editors, *Computer Aided Verification*, pages 443–452, Cham, 2019. Springer International Publishing.

[5] Mykel J Kochenderfer and JP Chryssanthacopoulos. Robust airborne collision avoidance through dynamic programming. *Massachusetts Institute of Technology, Lincoln Laboratory, Project Report ATC-371*, 130, 2011.

[6] Zhouxing Shi, Qirui Jin, J Zico Kolter, Suman Jana, Cho-Jui Hsieh, and Huan Zhang. Formal verification for neural networks with general nonlinearities via branch-and-bound. *2nd Workshop on Formal Verification of Machine Learning (WFVML 2023)*, 2023.

[7] Shiqi Wang, Huan Zhang, Kaidi Xu, Xue Lin, Suman Jana, Cho-Jui Hsieh, and J Zico Kolter. Beta-CROWN: Efficient bound propagation with per-neuron split constraints for complete and incomplete neural network verification. *Advances in Neural Information Processing Systems*, 34, 2021.

[8] Kaidi Xu, Zhouxing Shi, Huan Zhang, Yihan Wang, Kai-Wei Chang, Minlie Huang, Bhavya Kailkhura, Xue Lin, and Cho-Jui Hsieh. Automatic perturbation analysis for scalable certified robustness and beyond. *Advances in Neural Information Processing Systems*, 33, 2020.

[9] Kaidi Xu, Huan Zhang, Shiqi Wang, Yihan Wang, Suman Jana, Xue Lin, and Cho-Jui Hsieh. Fast and Complete: Enabling complete neural network verification with rapid and massively parallel incomplete verifiers. In *International Conference on Learning Representations*, 2021.

[10] Huan Zhang, Shiqi Wang, Kaidi Xu, Linyi Li, Bo Li, Suman Jana, Cho-Jui Hsieh, and J Zico Kolter. General cutting planes for bound-propagation-based neural network verification. *Advances in Neural Information Processing Systems*, 2022.

[11] Huan Zhang, Tsui-Wei Weng, Pin-Yu Chen, Cho-Jui Hsieh, and Luca Daniel. Efficient neural network robustness certification with general activation functions. *Advances in Neural Information Processing Systems*, 31:4939–4948, 2018.